

---

# Machine Learning Tricks

Philipp Koehn



- Myth of machine learning
  - given: real world examples
  - automatically build model
  - make predictions■
- Promise of deep learning
  - do not worry about specific properties of problem
  - deep learning automatically discovers the feature■
- Reality: bag of tricks

# Today's Agenda



- No new translation model
- Discussion of failures in machine learning
- Various tricks to address them

# Fair Warning



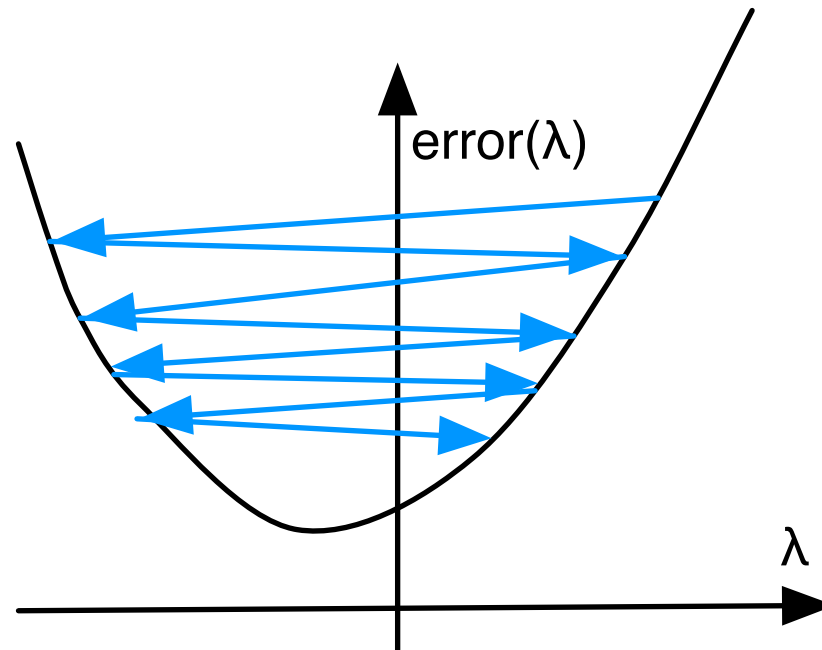
- At some point, you will think:

*Why are you telling us all this madness?*

- Because pretty much all of it is commonly used

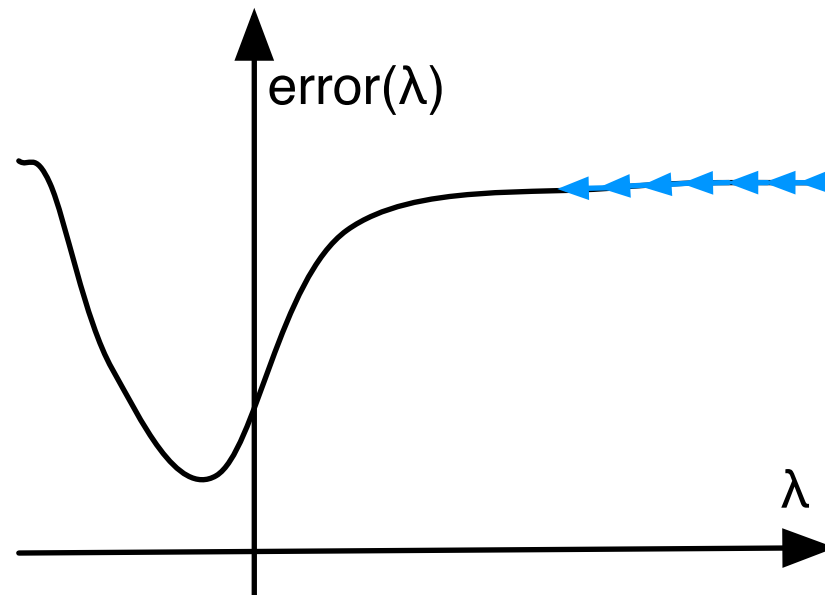
# failures in machine learning

# Failures in Machine Learning



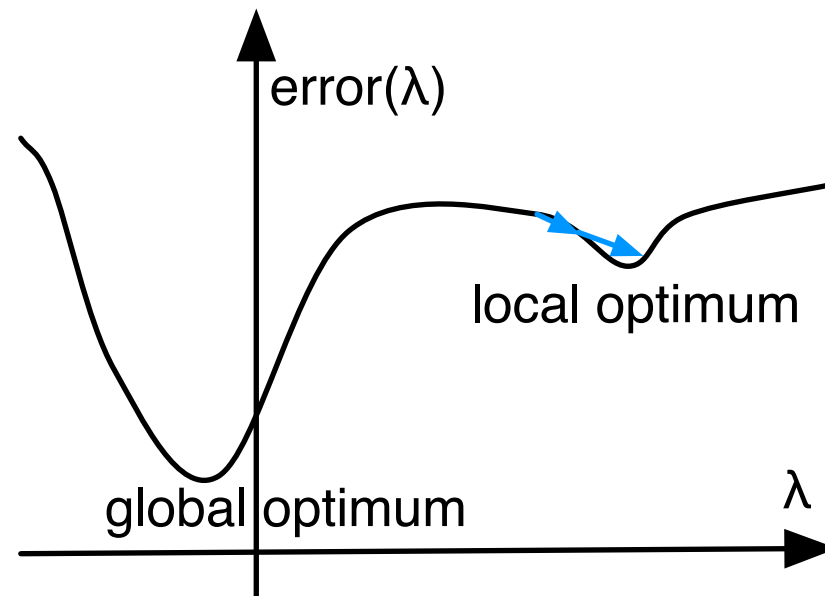
Too high learning rate may lead to too drastic parameter updates  
→ overshooting the optimum

# Failures in Machine Learning



Bad initialization may require many updates to escape a plateau

# Failures in Machine Learning



Local optima trap training



# Learning Rate



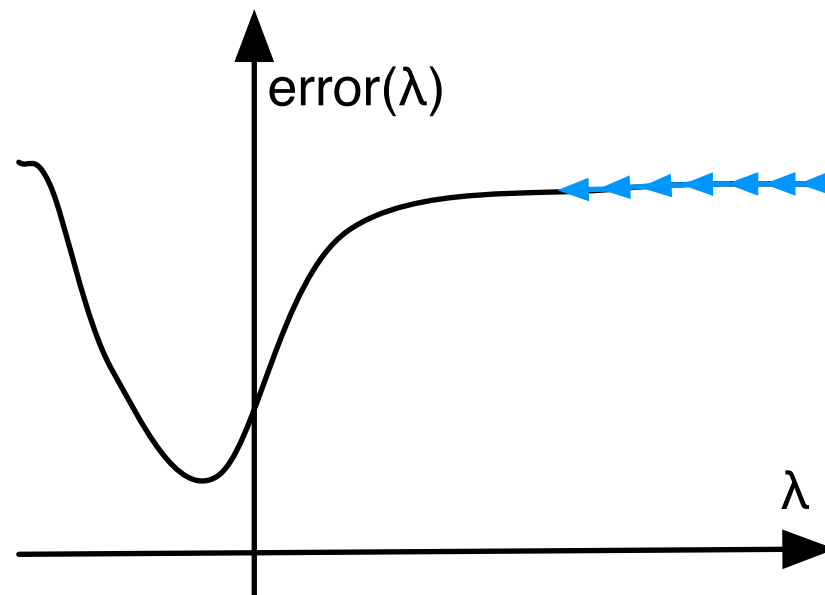
- Gradient computation gives direction of change
- Scaled by learning rate
- Weight updates■
- Simplest form: fixed value■
- Annealing
  - start with larger value (big changes at beginning)
  - reduce over time (minor adjustments to refine model)

# Initialization of Weights

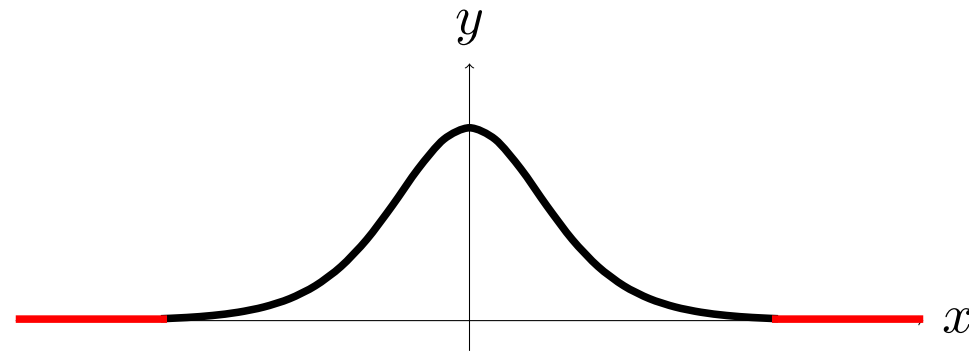


9

- Initialize weights to random values
- But: range of possible values matters



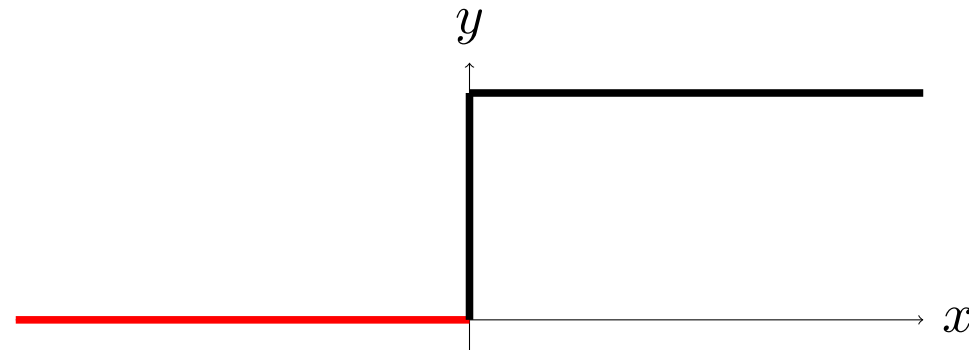
# Sigmoid Activation Function



Derivative of sigmoid

Near zero for large positive and negative values

# Rectified Linear Unit



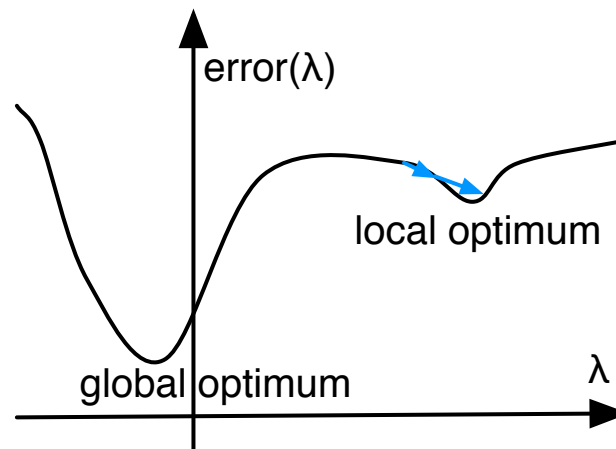
## Derivative of ReLU

Flat and for large interval: Gradient is 0

“Dead cells” elements in output that are always 0, no matter the input

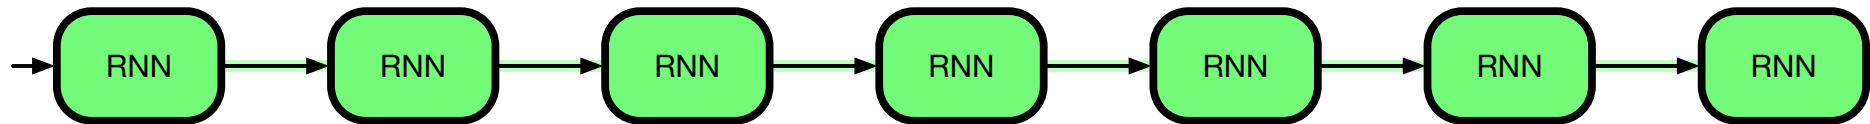
# Local Optima

- Cartoon depiction



- Reality
  - highly dimensional space
  - complex interaction between individual parameter changes
  - "bumpy"

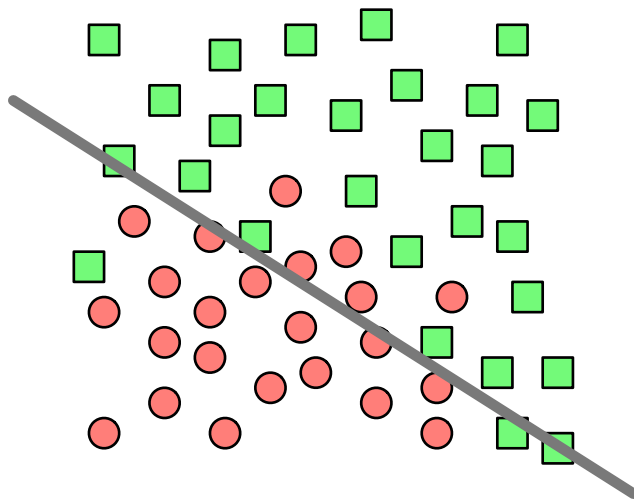
# Vanishing and Exploding Gradients



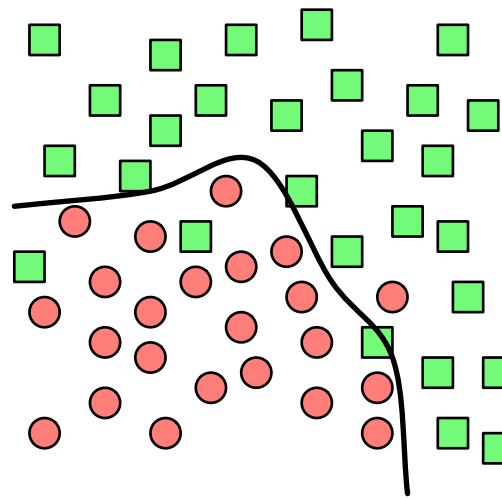
- Repeated multiplication with same values
- If gradients are too low  $\rightarrow 0$
- If gradients are too big  $\rightarrow \infty$

# Overfitting and Underfitting

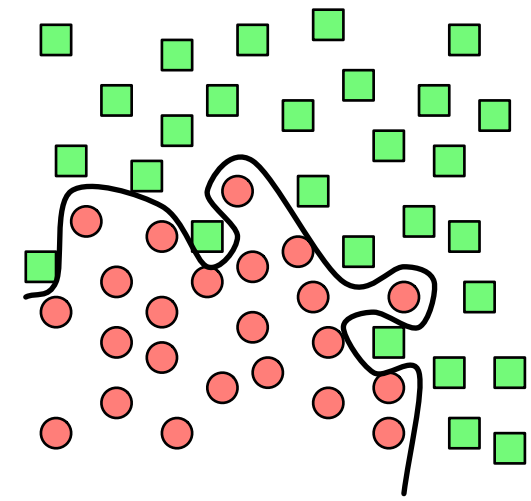
Under-Fitting



Good Fit



Over-Fitting



- Complexity of the problem has too match the capacity of the model
- Capacity  $\simeq$  number of trainable parameters

# ensuring randomness



# Ensuring Randomness

- Typical theoretical assumption

*independent and identically distributed*

training examples■

- Approximate this ideal
  - avoid undue structure in the training data
  - avoid undue structure in initial weight setting■
- ML approach: Maximum entropy training
  - Fit properties of training data
  - Otherwise, model should be as random as possible (i.e., has maximum entropy)

# Shuffling the Training Data

- Typical training data in machine translation
    - different types of corpora
      - \* European Parliament Proceedings
      - \* collection of movie subtitles
    - temporal structure in each corpus
    - similar sentences next too each other (e.g., same story / debate)■
  - Online updating: last examples matter more■
  - Convergence criterion: no improvement recently
    - stretch of hard examples following easy examples: prematurely stopped■
- ⇒ randomly shuffle the training data  
(maybe each epoch)

# Weight Initialization

- Initialize weights to random values
- Values are chosen from a uniform distribution
- Ideal weights lead to node values in transition area for activation function

## For Example: Sigmoid

- Input values in range  $[-1; 1]$

⇒ Output values in range  $[0.269; 0.731]$  ■

- Magic formula ( $n$  size of the previous layer)

$$\left[ -\frac{1}{\sqrt{n}}, \frac{1}{\sqrt{n}} \right] \blacksquare$$

- Magic formula for hidden layers

$$\left[ -\frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}}, \frac{\sqrt{6}}{\sqrt{n_j + n_{j+1}}} \right]$$

- $n_j$  is the size of the previous layer
- $n_{j+1}$  size of next layer

# Problem: Overconfident Models

- Predictions of the neural machine translation models are surprisingly confident
- Often almost all the probability mass is assigned to a single word (word prediction probabilities of over 99%)■
- Problem for decoding and training
  - decoding: sensible alternatives get low scores, bad for beam search
  - training: overfitting is more likely■
- Solution: label smoothing■
- Jargon notice
  - in classification tasks, we predict a *label*
  - jargon term for any output
  - here, we smooth the word predictions

# Label Smoothing during Decoding

- Common strategy to combat peaked distributions: smooth them
- Recall
  - prediction layer produces numbers for each word
  - converted into probabilities using the softmax

$$p(y_i) = \frac{\exp s_i}{\sum_j \exp s_j}$$

- Softmax calculation can be smoothed with so-called **temperature**  $T$

$$p(y_i) = \frac{\exp s_i/T}{\sum_j \exp s_j/T}$$

- Higher temperature  $\rightarrow$  distribution smoother  
(i.e., less probability is given to most likely choice)

# Label Smoothing during Training

- Root of problem: training
- Training object: assign all probability mass to single correct word
- Label smoothing
  - truth gives some probability mass to other words (say, 10% of it)
  - uniformly distributed over all words
  - relative to unigram word probabilities  
(relative counts of each word in the target side of the training data)

# adjusting the learning rate



# Adjusting the Learning Rate

- Gradient descent training: weight update follows the gradient downhill
- Actual gradients have fairly large values, scale with a learning rate (low number, e.g.,  $\mu = 0.001$ )
- Change the learning rate over time
  - starting with larger updates
  - refining weights with smaller updates
  - adjust for other reasons
- Learning rate schedule

# Momentum Term

- Consider case where weight value far from optimum
- Most training examples push the weight value in the same direction
- Small updates take long to accumulate
- Solution: momentum term  $m_t$ 
  - accumulate weight updates at each time step  $t$
  - some decay rate for sum (e.g., 0.9)
  - combine momentum term  $m_{t-1}$  with weight update value  $\Delta w_t$

$$m_t = 0.9m_{t-1} + \Delta w_t$$

$$w_t = w_{t-1} - \mu m_t$$

# Adapting Learning Rate per Parameter

- Common strategy: reduce the learning rate  $\mu$  over time
- Initially parameters are far away from optimum  $\rightarrow$  change a lot
- Later nuanced refinements needed  $\rightarrow$  change little
- Now: different learning rate for each parameter

# Adagrad

- Different parameters at different stages of training  
→ different learning rate for each parameter
- Adagrad
  - record gradients for each parameter
  - accumulate their square values over time
  - use this sum to reduce learning rate■
- Update formula
  - gradient  $g_t = \frac{dE_t}{dw}$  of error  $E$  with respect to weight  $w$
  - divide the learning rate  $\mu$  by accumulated sum

$$\Delta w_t = \frac{\mu}{\sqrt{\sum_{\tau=1}^t g_{\tau}^2}} g_t$$

- Big changes in the parameter value (corresponding to big gradients  $g_t$ )  
→ reduction of the learning rate of the weight parameter

# Adam: Elements

- Combine idea of momentum term and reduce parameter update by accumulated change
- Momentum term idea (e.g.,  $\beta_1 = 0.9$ )

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

- Accumulated gradients (decay with  $\beta_2 = 0.999$ )

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

# Adam: Technical Correction

- Initially, values for  $m_t$  and  $v_t$  are close to initial value of 0
- Adjustment

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}, \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- With  $t \rightarrow \infty$  this correction goes away

$$\lim_{t \rightarrow \infty} \frac{1}{1 - \beta^t} \rightarrow 1$$

- Given
  - learning rate  $\mu$
  - momentum  $\hat{m}_t$
  - accumulated change  $\hat{v}_t$
- Weight update per Adam (e.g.,  $\epsilon = 10^{-8}$ )

$$\Delta w_t = \frac{\mu}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$

# Batched Gradient Updates

- Accumulate all weight updates for all the training example → update (converges slowly)■
- Process each training example → update (stochastic gradient descent) (quicker convergence, but last training disproportionately higher impact)■
- Process data in batches
  - compute all their gradients for individual word predictions errors
  - use sum over each batch to update parameters→ better parallelization on GPUs■
- Process data on multiple compute cores
  - batch processing may take different amount of time
  - asynchronous training: apply updates when they arrive
  - mismatch between original weights and updates may not matter much



# avoiding local optima

# Avoiding Local Optima

- One of the hardest problem for designing neural network architectures and optimization methods
- Ensure that model converges to at least to a set of parameter values that give results close to this optimum on unseen test data.
- There is no real solution to this problem.
- It requires experimentation and analysis that is more craft than science.
- Still, this section presents a number of methods that generally help avoiding getting stuck in local optima.

# Overfitting and Underfitting

- Neural machine translation models
  - 100s of millions of parameters
  - 100s of millions of training examples (individual word predictions)
- No hard rules for relationship between these two numbers
- Too many parameters and too few training examples → overfitting
- Too few parameters and many training examples → underfitting

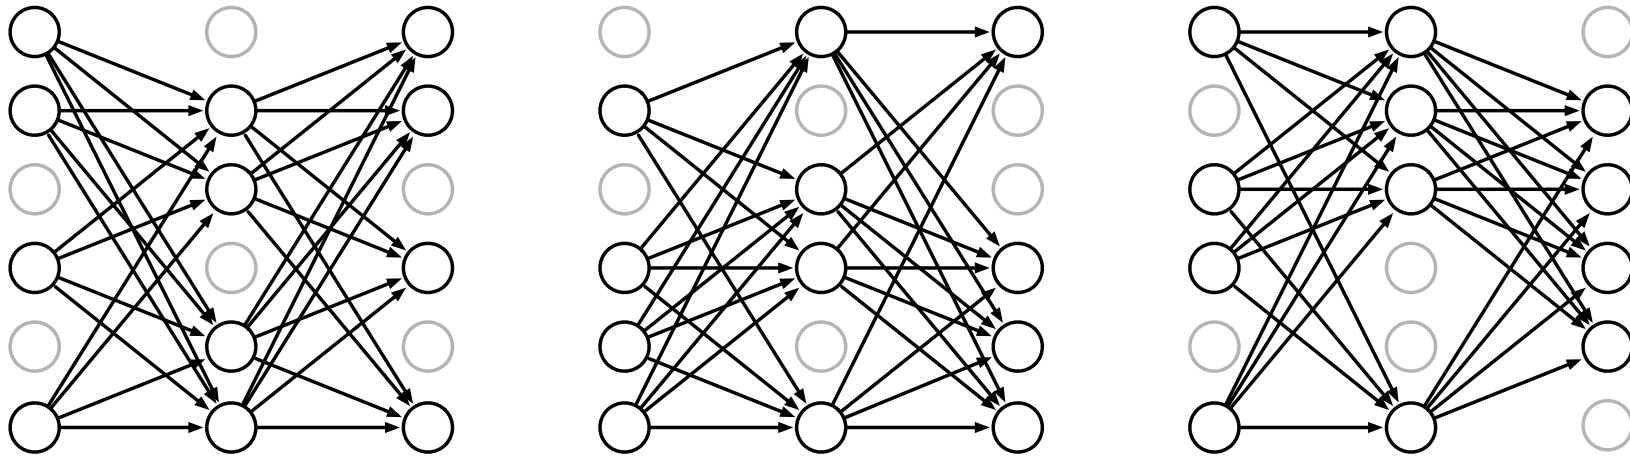
# Regularization

- Motivation: prefer as few parameters as possible
- Strategy: set un-needed parameters a value of 0
- Method
  - adjust training objective
  - add cost for any non-zero parameter
  - typically done with L2 norm
- Practical impact
  - derivative of L2 norm is value of parameter
  - if not signal from training: reduce value of parameter
  - also called weight decay
- Not common in deep learning, but other methods understood as regularization

- Human learning
  - learn simple concepts first
  - learn more complex material later
- Early epochs: only easy training examples
  - only short sentences
  - create artificial data by extracting smaller segments  
(similar to phrase pair extraction in statistical machine translation)
  - Later epochs: all training data
- Not easy to calibrate

- Training may get stuck in local optima
  - some properties of task have been learned
  - discovery of other properties would take it too far out of its comfort zone.
- Machine translation example
  - model learned the language model aspects
  - but cannot figure out role of input sentence
- Drop out: for each batch, eliminate some nodes

# Dropout



- Dropout
  - For each batch, different random set of nodes is removed
  - Their values are set to 0 and their weights are not updated
  - 10%, 20% or even 50% of all the nodes
- Why does this work?
  - robustness: redundant nodes play similar nodes
  - ensemble learning: different subnetworks are different models

# Gradient Clipping

- Exploding gradients: gradients become too large during backward pass

⇒ Limit total value of gradients for a layer to threshold ( $\tau$ )

- Use of L2 norm of gradient values  $g$

$$L2(g) = \sqrt{\sum_j g_j^2}$$

- Adjust each gradient value  $g_i$  for each element  $i$  in the vector

$$g'_i = g_i \times \frac{\tau}{\max(\tau, L2(g))}$$



# Layer Normalization

- During inference, average node values may become too large or too small
- Has also impact on training (gradients are multiplied with node values)

⇒ Normalize node values

- During training, learn bias layer

# Layer Normalization: Math

- Feed-forward layer  $h^l$ , weights  $W$ , computed sum  $s^l$ , activation function

$$s^l = W h^{l-1}$$

$$h^l = \text{sigmoid}(s^l)$$

- Compute mean  $\mu^l$  and variance  $\sigma^l$  of sum vector  $s^l$

$$\mu^l = \frac{1}{H} \sum_{i=1}^H s_i^l$$

$$\sigma^l = \sqrt{\frac{1}{H} \sum_{i=1}^H (s_i^l - \mu^l)^2}$$

# Layer Normalization: Math

- Normalize  $s^l$

$$\hat{s}^l = \frac{1}{\sigma^l}(s^l - \mu^l)$$

- Learnable bias vectors  $g$  and  $b$

$$\hat{s}^l = \frac{g}{\sigma^l}(s^l - \mu^l) + b$$

# Shortcuts and Highways

- Deep learning: many layers of processing

⇒ Error propagation has to travel farther

- All parameters in processing change have to be adjusted
- Instead of always passing through all layers, add connections from first to last
- Jargon alert
  - shortcuts
  - residual connections
  - skip connections

- Feed-forward layer

$$y = f(x)$$

- Pass through input  $x$

$$y = f(x) + x$$

- Note: gradient is

$$y' = f'(x) + 1$$

- Constant 1  $\rightarrow$  gradient is passed through unchanged

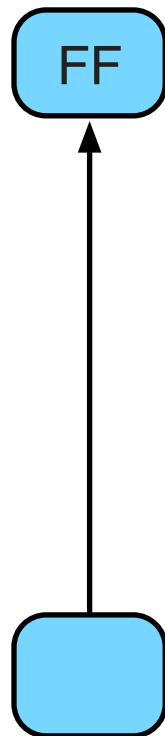
# Highways

- Regulate how much information from  $f(x)$  and  $x$  should impact the output  $y$
- Gate  $t(x)$  (typically computed by a feed-forward layer)

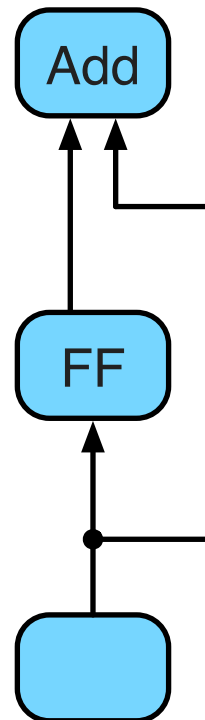
$$y = t(x) f(x) + (1 - t(x)) x$$

# Shortcuts and Highways


Basic Layer



Skip Connection



Highway Network



# LSTM and Vanishing Gradients

- Recall: Long short term memory (LSTM) cells
- Pass through of memory

$$\text{memory}^t = \text{gate}_{\text{input}} \times \text{input}^t + \text{gate}_{\text{forget}} \times \text{memory}^{t-1}$$

- Forget gate has values close to 1  $\rightarrow$  gradient passed through nearly unchanged



# generative adversarial training

# Sequence-Level Training

- Traditional training
  - predict one word at a time
  - compare against correct word
  - proceed training with correct word
- Sequence-level training
  - predict entire sequence
  - measure translation with sentence-level metric (e.g., BLEU)
- May use n-best translations, beam search, etc.

- Game between two players
  - generator proposes a translation
  - discriminator distinguishes between generator's translation and human translation
  - generator tries to fool discriminator
- Training example: input sentence  $x$  and output sentence  $y$
- Generator
  - traditional neural machine translation model
  - generates full sentence translations  $t$  for each input sentence
- Discriminator
  - is trained to classify  $(x, y)$  as correct example
  - is trained to classify  $(x, t)$  as generated example



1. First train generator to some maturity
  2. Train discriminator on generator predictions and human reference translations
  3. Train jointly
    - generator with additional objective to fool discriminator
    - discriminator to do well on detecting generator's output as such
- In practice, this is hard to calibrate correctly

# Relationship to Reinforcement Learning

- No immediate feedback
  - chess playing: quality of move only revealed at end of game
  - walk through maze to avoid monsters and find gold
- Policy: decision process to which steps to take  
(here: generator, traditional neural machine translation model)
- Reward: end result  
(here: ability to fool discriminator)
- Popular technique: Monte Carlo search  
(here: Monte Carlo decoding)
- Training is called policy search