
Alternative Architectures

Philipp Koehn



Alternative Architectures



- We introduced one translation model
 - attentional seq2seq model
 - core organizing feature: recurrent neural networks
- Other core neural architectures
 - convolutional neural networks
 - attention
- But first: look at various components of neural architectures

components

Components of Neural Networks



- Neural networks originally inspired by the brain
 - a neuron receives signals from other neurons
 - if sufficiently activated, it sends signals
 - feed-forward layers are roughly based on this
- Computation graph
 - any function possible, as long as it is partially differentiable
 - not limited by appeals to biological validity
- *Deep learning* maybe a better name

Feed-Forward Layer



- Classic neural network component
- Given an input vector x , matrix multiplication M with adding a bias vector b

$$Mx + b$$

- Adding a non-linear activation function

$$y = \text{activation}(Mx + b)$$

- Notation

$$y = FF_{\text{activation}}(x) = a(Mx + b)$$

Feed-Forward Layer



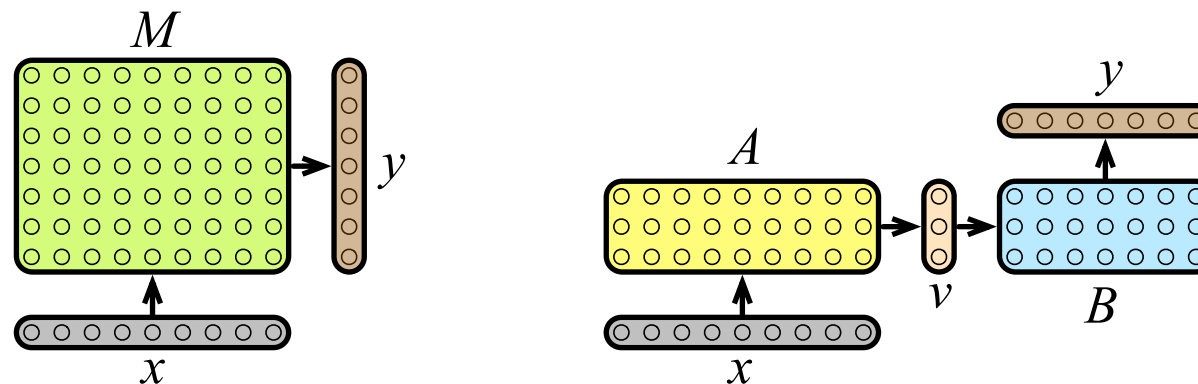
- Historic neural network designs: several feed-forward layers
 - input layer
 - hidden layers
 - output layer
- Powerful tools for a wide range of machine learning problems
- Matrix multiplication also called **affine transforms**
 - appeals to its geometrical properties
 - straight lines in input still straight lines in output

Factored Decomposition


- One challenge: very large input and output vectors
- Number of parameters in matrix $M = |x| \times |y|$

⇒ Need to reduce size of matrix

- Solution: first reduce to smaller representation



Factored Decomposition: Math



- Intuition

- given highly dimension vector x
- first map to into lower dimensional vector v (matrix A)
- then map to output vector y (matrix B)

$$v = Ax$$

$$y = Bv = BAx$$

- Example

- $|x| = 20,000, |y| = 50,000 \rightarrow M = 1,000,000,000$
- $|v| = 100 \rightarrow A = 20,000 \times 100 = 2,000,000, B = 100 \times 50,000 = 5,000,000$
- reduction from 1,000,000,000 to 7,000,000

Factored Decomposition: Interpretation



- Vector v is a bottleneck feature
- Forced to captures salient features
- One example: word embeddings



processing sequences

- Already described recurrent neural networks at length
 - propagate state s
 - over time steps t
 - receiving an input x_t at each turn

$$s_t = f(s_{t-1}, x_t)$$

(state may computed may as a feed-forward layer)■

- More successful
 - gated recurrent units (GRU)
 - long short-term memory cells (LSTM)■
- Good fit for sequences, like words in a sentence
 - humans also receive word by word
 - most recent words most relevant
 - closer to current state
- But computational problematic: very long computation chains

Alternative Sequence Processing




- Convolutional neural networks
- Attention




convolutional neural networks

Convolutional Neural Networks (CNN)



- Popular in image processing
- Regions of an image are reduced into increasingly smaller representation
 - matrix spanning part of image reduced to single value
 - overlapping regions



- Map words into fixed-sized sentence representation

Hierarchical Structure and Language

- Syntactic and semantic theories of language
 - language is recursive
 - central: verb
 - dependents: subject, objects, adjuncts
 - their dependents: adjectives, determiners
 - also nested: relative clauses
- How to compute sentence embeddings active research topic

- Key step
 - take a high dimensional input representation
 - map to lower dimensional representation
- Several repetitions of this step■
- Examples
 - map 50×50 pixel area into scalar value
 - combine 3 or more neighboring words into a single vector■
- Machine translation
 - encode input sentence into single vector
 - decode this vector into a sentence in the output language

attention

- Machine translation is a structured prediction task
 - output is not a single label
 - output structure needs to be built, word by word
- Relevant information for each word prediction varies
- Human translators pay attention to different parts of the input sentence when translating

⇒ Attention mechanism

- Attention mechanism in neural translation model (Bahdanau et al., 2015)
 - previous hidden state s_{i-1}
 - input word embedding h_j
 - trainable parameters b, W_a, U_a, v_a

$$a(s_{i-1}, h_j) = v_a^T \tanh(W_a s_{i-1} + U_a h_j + b)$$

- Other ways to compute attention
 - Dot product: $a(s_{i-1}, h_j) = s_{i-1}^T h_j$
 - Scaled dot product: $a(s_{i-1}, h_j) = \frac{1}{\sqrt{|h_j|}} s_{i-1}^T h_j$
 - General: $a(s_{i-1}, h_j) = s_{i-1}^T W_a h_j$
 - Local: $a(s_{i-1}) = W_a s_{i-1}$

Attention of Luong et al. (2015)


- Luong et al. (2015) demonstrate good results with the dot product

$$a(s_{i-1}, h_j) = s_{i-1}^T h_j$$


- No trainable parameters
- Additional changes
- Currently more popular

Attention of Luong et al. (2015)

Luong et al. (2015)



Bahdanau et al. (2015)



Attention of Luong et al. (2015)

Luong et al. (2015)

Attention

$$\alpha_{ij} = \text{softmax FF}(s_{i-1}, h_j)$$

Input context $c_i = \sum_j \alpha_{ij} h_j$

Output word

$$p(y_t | y_{<t}, x) = \text{softmax}(W \text{ FF}_{\tanh}(s_{i-1}, c_i))$$

Decoder state

$$s_i = \text{FF}_{\tanh}(s_{i-1}, Ey_{i-1})$$

Bahdanau et al. (2015)

Attention

$$\alpha_{ij} = \text{softmax FF}(s_{i-1}, h_j)$$

Input context $c_i = \sum_j \alpha_{ij} h_j$

Output word

$$p(y_t | y_{<t}, x) = \text{softmax}(W \text{ FF}_{\tanh}(s_{i-1}, Ey_{i-1}, c_i))$$

Decoder state

$$s_i = \text{FF}_{\tanh}(s_{i-1}, Ey_{i-1}, c_i)$$

Multi-Head Attention

- Add redundancy
 - say, 16 attention weights
 - each based on its own parameters■
- Formally, for each head k compute an associated between
 - decoder state s_{i-1} at time step i
 - encoder state h_j for the j th input word
 - using the softmax of some parameterized function a^k

$$\alpha_{ij}^k = \text{softmax } a^k(s_{i-1}, h_j) \blacksquare$$

- Average the attention weights

$$\alpha_{ij} = \frac{1}{k} \sum_k \alpha_{ij}^k$$

- Multi-head attention is a form of ensembling

Fine-Grained Attention

- Why just use a single scalar value to weight entire vectors?
 - learn weights for each element
 - computation of attention values returns vector instead of scalar■
- Architecturally, still a feed-forward neural network (or any of variants)

$$a(s_{i-1}, h_j) = \text{FF}^k(s_{i-1}, h_j) \blacksquare$$

- Softmax is now applied over each dimension d

$$\alpha_{ij}^d = \frac{\exp a^d(s_{i-1}, h_j)}{\sum_k a^d(s_{i-1}, h_k)} \blacksquare$$

- Input context is now computed by a element-wise multiplication

$$c_i = \sum_j \alpha_{ij} \times h_j$$

- Finally, a very different take at attention
- Motivation so far: need for alignment between input words and output words
- Now: refine representation of input words in the encoder
 - representation of an input word mostly depends on itself
 - but also informed by the surrounding context
 - previously: recurrent neural networks (considers left or right context)
 - now: attention mechanism
- Self attention:
Which of the surrounding words is most relevant to refine representation?

- Formal definition (based on sequence of vectors h_j , packed into matrix H)

$$\text{self-attention}(H) = \text{softmax}\left(\frac{H H^T}{\sqrt{|h|}}\right) H$$

- Association between every word representation h_j any other context word h_k
 - computed by dot product
 - results in a vector of raw association values

$$H H^T$$

- Scaled by the size of the word representation vectors $|h|$, and softmax

$$\text{softmax}\left(\frac{H H^T}{\sqrt{|h|}}\right)$$

- Resulting vector of normalized association values used to weigh context words

Self Attention

- More familiar math, using word representation vectors h_j

- Raw association $\frac{HH^T}{\sqrt{|h|}}$

$$a_{jk} = \frac{1}{|h|} h_j h_k^T$$

- Normalized association (softmax)

$$\alpha_{jk} = \frac{\exp(a_{jk})}{\sum_{\kappa} \exp(a_{j\kappa})}$$

- Weighted sum


$$\text{self-attention}(h_j) = \sum_k \alpha_{jk} h_k$$

- More on this later (Transformer)

convolutional machine translation

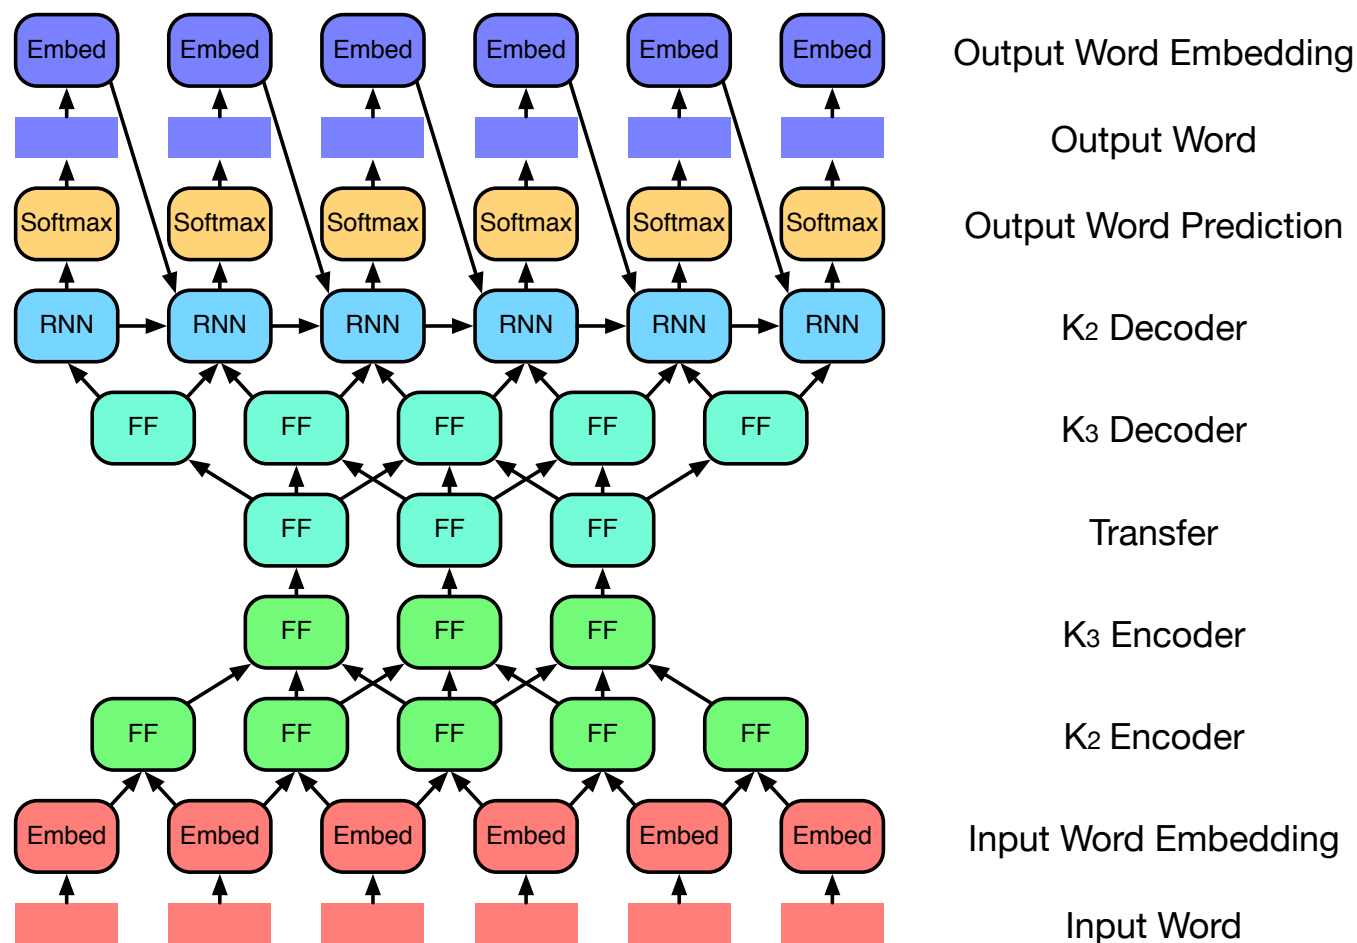
Convolutional Machine Translation

- First end-to-end neural machine translation model of the modern era
[Kalchbrenner and Blunsom, 2013]
- Encoder



- always two convolutional layers, with different size
 - here: K_2 and K_3
- Decoder similar

Refinement




- Convolutions do not result in a single sentence embedding but a sequence
- Decoder is also informed by a recurrent neural network

CNNs With Attention

[Gehring et al. 2017]

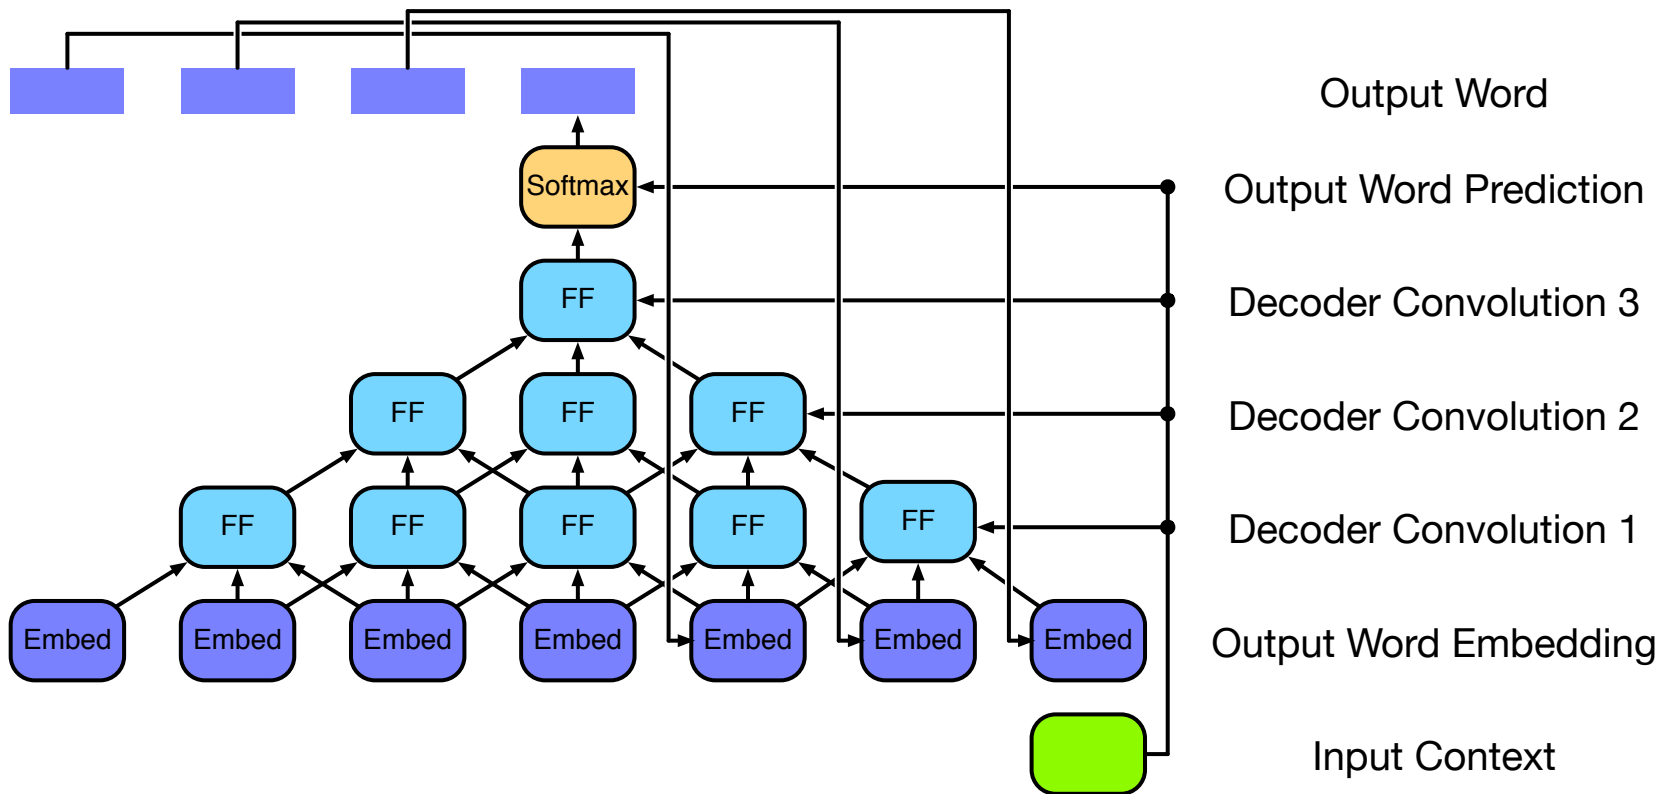
- Combination of
 - convolutional neural networks
 - attention
- Sequence-to-sequence attention, mainly as before
- Recurrent neural networks replaced by convolutional layers

Encoder



- Stacked encoder convolutions
- Not shortening representations
- But: faster processing due to more parallelism

Decoder




- Decoder state computed by convolutional layers over previous output words
- Each convolutional state also informed by the input context (using attention)

transformer

- Self-attention in encoder
 - refine word representation based on relevant context words
 - relevance determined by self attention
- Self-attention in decoder
 - refine output word predictions based on relevant previous output words
 - relevance determined by self attention
- Also regular attention to encoder states in decoder
- Currently most successful model
(maybe only with self attention in decoder, but regular recurrent decoder)

Encoder



Sequence of self-attention layers

Self Attention Layer

- Given: input word representations h_j , packed into a matrix $H = (h_1, \dots, h_j)$

- Self attention

$$\text{self-attention}(H) = \text{softmax}\left(\frac{HH^T}{\sqrt{|h|}}\right)H$$

- Shortcut connection

$$\text{self-attention}(h_j) + h_j$$

- Layer normalization

$$\hat{h}_j = \text{layer-normalization}(\text{self-attention}(h_j) + h_j)$$

- Feed-forward step with ReLU activation function

$$\text{relu}(W\hat{h}_j + b)$$

- Again, shortcut connection and layer normalization

$$\text{layer-normalization}(\text{relu}(W\hat{h}_j + b) + \hat{h}_j)$$

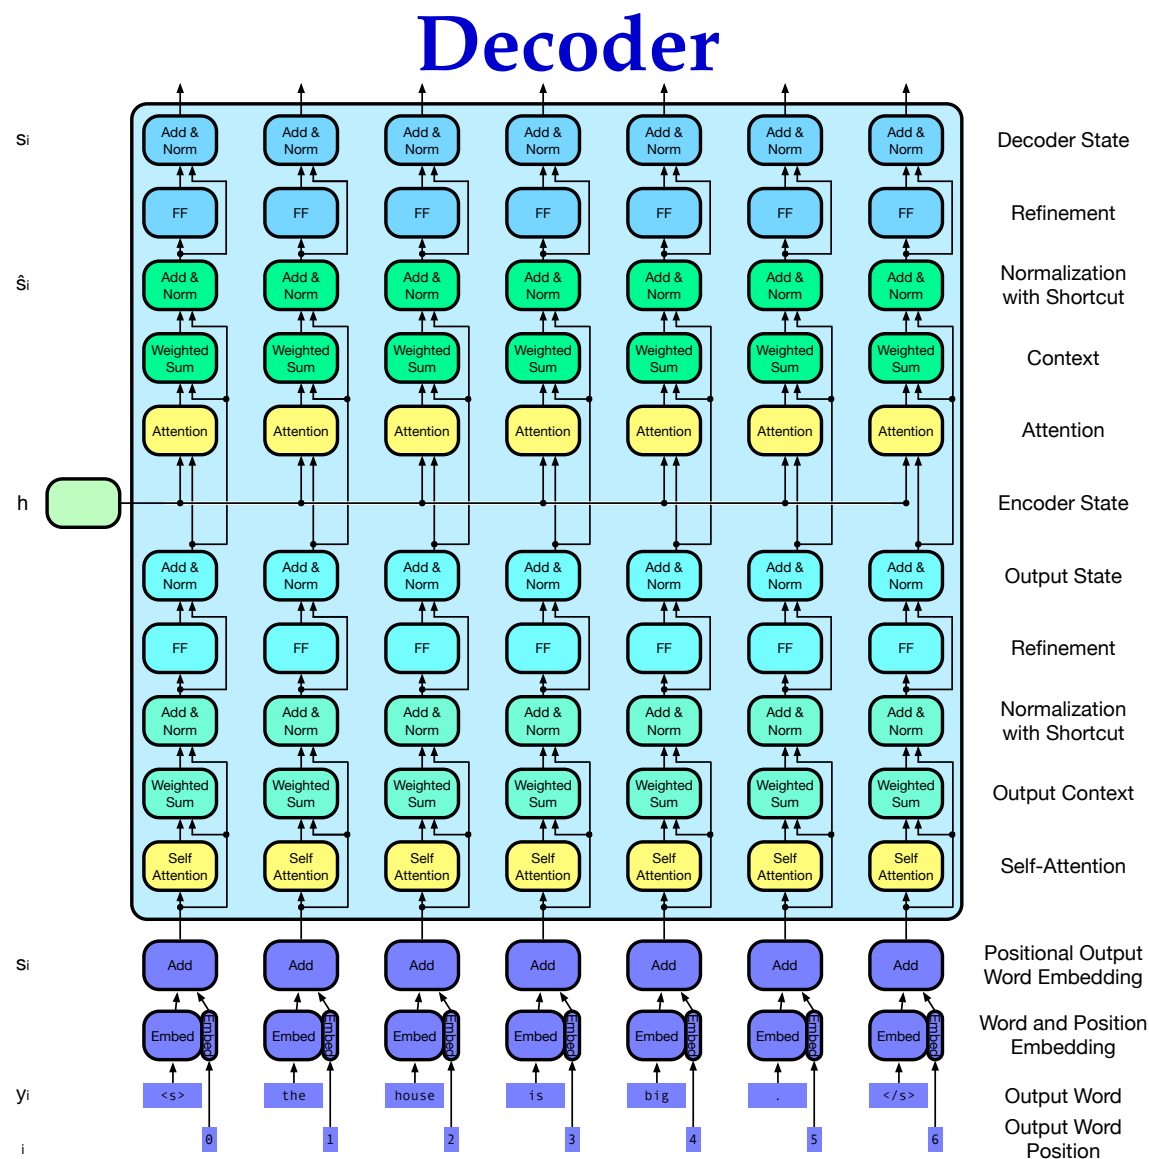
Stacked Self Attention Layers

- Stack several such layers (say, $D = 6$)
- Start with input word embedding

$$h_{0,j} = Ex_j$$

- Stacked layers

$$h_{d,j} = \text{self-attention-layer}(h_{d-1,j})$$



Decoder computes attention-based representations of the output in several layers, initialized with the embeddings of the previous output words

Self-Attention in the Decoder

- Same idea as in the encoder
- Output words are initially encoded by word embeddings $s_i = Ey_i$.
- Self attention is computed over previous output words
 - association of a word s_i is limited to words s_k ($k \leq i$)
 - resulting representation \tilde{s}_i

$$\text{self-attention}(\tilde{S}) = \text{softmax}\left(\frac{SS^T}{\sqrt{|h|}}\right)S$$

- Original intuition of attention mechanism: focus on relevant input words
- Computed with dot product $\tilde{S}H^T$
- Compute attention between the decoder states \tilde{S} and the final encoder states H

$$\text{attention}(\tilde{S}, H) = \text{softmax}\left(\frac{\tilde{S}H^T}{\sqrt{|h|}}\right)H$$

- Note: attention mechanism formally mirrors self-attention

Full Decoder

